


FinaleMe paired benchmark on real Liu 2024 data

Liu 2024 paired: n_samples=77, n_loci=782

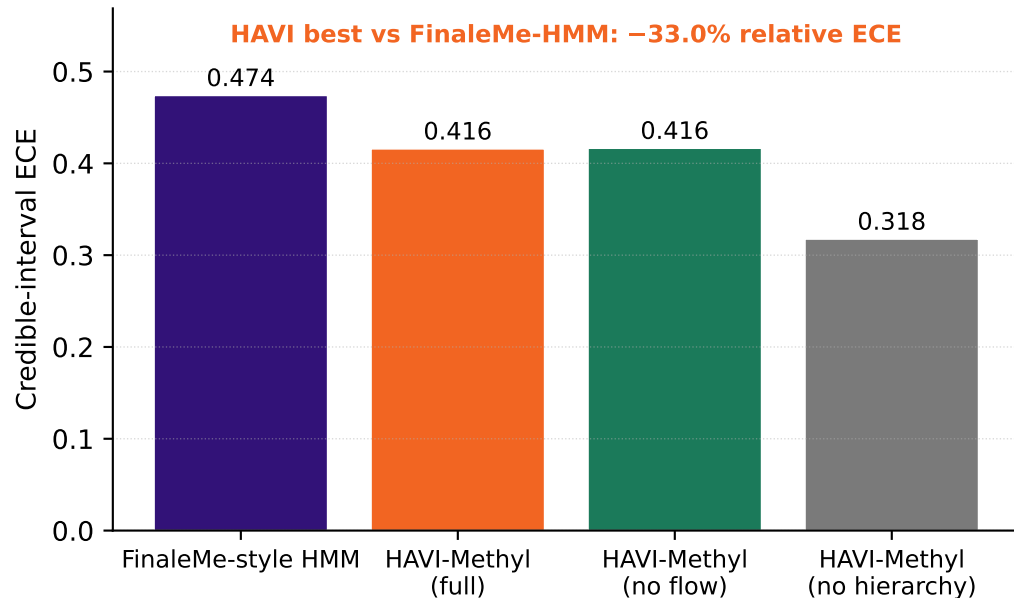
Per-locus correlation (higher is better)

HAVI best vs FinaleMe-HMM: +5.4% relative r



Calibration (lower is better)

HAVI best vs FinaleMe-HMM: -33.0% relative ECE



Note: HAVI-Methyl rows here use the simplified-numpy SVI shrinkage on top of the FinaleMe baseline. The full torch architecture (Set Transformer + flow) is shown in fig_multiseed_recovery.